


# Representer Theorem



# Purpose

In this lecture we discuss:

- The kernel trick
- The representer theorem
- Applications of the representer theorem
  - Surface reconstruction
  - Smoothing splines

# Representer Theorem

Recall the supervised learning setting:

- We are given training data  $\tau = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$  and a loss function that measures the fit to the data.
- We wish to find a function  $g \in \mathcal{G}$  that minimizes the training loss, with the addition of a regularization term.
- We assume that the class  $\mathcal{G}$  of prediction functions can be decomposed as the direct sum of an RKHS  $\mathcal{H}$ , defined by a kernel function  $\kappa : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ , and another linear space of real-valued functions  $\mathcal{H}_0$  on  $\mathcal{X}$ ; that is,

$$\mathcal{G} = \mathcal{H} \oplus \mathcal{H}_0,$$

meaning that any element  $g \in \mathcal{G}$  can be written as  $g = h + h_0$ , with  $h \in \mathcal{H}$  and  $h_0 \in \mathcal{H}_0$ . In minimizing the training loss we wish to penalize the  $h$  term of  $g$  but not the  $h_0$  term.

# Functional Optimization

The aim is to solve the functional optimization problem

$$\min_{g \in \mathcal{H} \oplus \mathcal{H}_0} \frac{1}{n} \sum_{i=1}^n \text{Loss}(y_i, g(\mathbf{x}_i)) + \gamma \|g\|_{\mathcal{H}}^2, \quad (1)$$

where  $\|g\|_{\mathcal{H}}$  means  $\|h\|_{\mathcal{H}}$  if  $g = h + h_0$ .

In this way, we can view  $\mathcal{H}_0$  as the **null space** of the functional  $g \mapsto \|g\|_{\mathcal{H}}$ .

This null space may be empty, but typically has a small dimension  $m$ ; for example it could be the one-dimensional space of constant functions.

## Example: Ridge Regression (cont.)

Consider again the ridge regression optimization problem

$$\min_{g \in \mathcal{H} \oplus \mathcal{H}_0} \frac{1}{n} \sum_{i=1}^n (y_i - g(\tilde{\mathbf{x}}_i))^2 + \gamma \|g\|_{\mathcal{H}}^2,$$

where we have feature vectors  $\tilde{\mathbf{x}} = [1, \mathbf{x}^\top]^\top$  and  $\mathcal{G}$  consists of functions of the form  $g : \tilde{\mathbf{x}} \mapsto \beta_0 + \mathbf{x}^\top \boldsymbol{\beta}$ .

Each function  $g$  can be decomposed as  $g = h + h_0$ , where  $h : \tilde{\mathbf{x}} \mapsto \mathbf{x}^\top \boldsymbol{\beta}$ , and  $h_0 : \tilde{\mathbf{x}} \mapsto \beta_0$ .

For  $g \in \mathcal{G}$ , we have  $\|g\|_{\mathcal{H}} = \|\boldsymbol{\beta}\|$ , and so the null space  $\mathcal{H}_0$  of the functional  $g \mapsto \|g\|_{\mathcal{H}}$  is the set of constant functions here, which has dimension  $m = 1$ .

## Example: Ridge Regression (cont.)

Regularization favors elements in  $\mathcal{H}_0$  and penalizes large elements in  $\mathcal{H}$ .

As the regularization parameter  $\gamma$  varies between zero and infinity, solutions to (1) vary from “complex” ( $g \in \mathcal{H} \oplus \mathcal{H}_0$ ) to “simple” ( $g \in \mathcal{H}_0$ ).

By choosing  $\mathcal{H}$  to be an RKHS in (1) this **functional** optimization problem effectively becomes a **parametric** optimization problem.

The reason is that any solution to (1) can be represented as a finite-dimensional linear combination of kernel functions, evaluated at the training sample.

This is due to the next **representer theorem** and is known as the **kernel trick**.

## Theorem: Representer Theorem

Let  $\mathcal{H}$  be an RKHS with kernel  $\kappa$ . The solution to the penalized optimization problem

$$\min_{g \in \mathcal{H} \oplus \mathcal{H}_0} \frac{1}{n} \sum_{i=1}^n \text{Loss}(y_i, g(\mathbf{x}_i)) + \gamma \|g\|_{\mathcal{H}}^2,$$

is of the form

$$g(\mathbf{x}) = \sum_{i=1}^n \alpha_i \kappa(\mathbf{x}_i, \mathbf{x}) + \sum_{j=1}^m \eta_j q_j(\mathbf{x}), \quad (2)$$

where  $\{q_1, \dots, q_m\}$  is a basis of  $\mathcal{H}_0$ .

## Proof

Let  $\mathcal{F} = \text{Span} \{ \kappa_{\mathbf{x}_i}, i = 1, \dots, n \}$ . Clearly,  $\mathcal{F} \subseteq \mathcal{H}$ .

Let  $\mathcal{F}^\perp$  be the orthogonal complement of  $\mathcal{F}$ ; that is, the class of functions

$$\{ f^\perp \in \mathcal{H} : \langle f^\perp, f \rangle_{\mathcal{H}} = 0, f \in \mathcal{F} \} \equiv \{ f^\perp : \langle f^\perp, \kappa_{\mathbf{x}_i} \rangle_{\mathcal{H}} = 0, \forall i \}.$$

Then  $\mathcal{H} = \mathcal{F} \oplus \mathcal{F}^\perp$ .

By the reproducing kernel property, for all  $f^\perp \in \mathcal{F}^\perp$ :

$$f^\perp(\mathbf{x}_i) = \langle f^\perp, \kappa_{\mathbf{x}_i} \rangle_{\mathcal{H}} = 0, \quad i = 1, \dots, n.$$



## Proof (cont.)

Take any  $g \in \mathcal{H} \oplus \mathcal{H}_0$ , and write it as  $g = f + f^\perp + h_0$ , with  $f \in \mathcal{F}$ ,  $f^\perp \in \mathcal{F}^\perp$ , and  $h_0 \in \mathcal{H}_0$ .

By the definition of the null space  $\mathcal{H}_0$ , we have  $\|g\|_{\mathcal{H}}^2 = \|f + f^\perp\|_{\mathcal{H}}^2$ .

By Pythagoras,  $\|f + f^\perp\|_{\mathcal{H}}^2 = \|f\|_{\mathcal{H}}^2 + \|f^\perp\|_{\mathcal{H}}^2$ .

It follows that

$$\begin{aligned} \frac{1}{n} \sum_{i=1}^n \text{Loss}(y_i, g(\mathbf{x}_i)) + \gamma \|g\|_{\mathcal{H}}^2 &= \frac{1}{n} \sum_{i=1}^n \text{Loss}(y_i, f(\mathbf{x}_i) + h_0(\mathbf{x}_i)) + \gamma (\|f\|_{\mathcal{H}}^2 + \|f^\perp\|_{\mathcal{H}}^2) \\ &\geq \frac{1}{n} \sum_{i=1}^n \text{Loss}(y_i, f(\mathbf{x}_i) + h_0(\mathbf{x}_i)) + \gamma \|f\|_{\mathcal{H}}^2. \end{aligned}$$

Since we can obtain equality by taking  $f^\perp = 0$ , this implies that the minimizer of the penalized optimization problem (1) lies in the subspace  $\mathcal{F} \oplus \mathcal{H}_0$  of  $\mathcal{G} = \mathcal{H} \oplus \mathcal{H}_0$ , and hence is of the form (2).  $\square$

# Parametric Optimization

Substituting the representation (2) of  $g$  into (1) gives the **finite-dimensional** parametric optimization problem:

$$\min_{\alpha \in \mathbb{R}^n, \eta \in \mathbb{R}^m} \frac{1}{n} \sum_{i=1}^n \text{Loss}(y_i, (\mathbf{K}\alpha + \mathbf{Q}\eta)_i) + \gamma \alpha^\top \mathbf{K}\alpha, \quad (3)$$

where

- $\mathbf{K}$  is the  $n \times n$  (Gram) matrix with entries  $[\kappa(\mathbf{x}_i, \mathbf{x}_j), i = 1, \dots, n, j = 1, \dots, n]$ .
- $\mathbf{Q}$  is the  $n \times m$  matrix with entries  $[q_j(\mathbf{x}_i), i = 1, \dots, n, j = 1, \dots, m]$ .

# Convex Optimization

In particular, for the squared-error loss we have

$$\min_{\alpha \in \mathbb{R}^n, \eta \in \mathbb{R}^m} \frac{1}{n} \| \mathbf{y} - (\mathbf{K}\alpha + \mathbf{Q}\eta) \|^2 + \gamma \alpha^\top \mathbf{K}\alpha. \quad (4)$$

This is a convex optimization problem, and its solution is found by differentiating (4) with respect to  $\alpha$  and  $\eta$  and equating to zero, leading to the following system of  $(n + m)$  linear equations:

$$\begin{bmatrix} \mathbf{K}\mathbf{K}^\top + n\gamma\mathbf{K} & \mathbf{K}\mathbf{Q} \\ \mathbf{Q}^\top\mathbf{K}^\top & \mathbf{Q}^\top\mathbf{Q} \end{bmatrix} \begin{bmatrix} \alpha \\ \eta \end{bmatrix} = \begin{bmatrix} \mathbf{K}^\top \\ \mathbf{Q}^\top \end{bmatrix} \mathbf{y}. \quad (5)$$

As long as  $\mathbf{Q}$  is of full column rank, the minimizing function is unique.

## Example: Ridge Regression (cont.)

Recall the ridge regression minimization program:

$$\min_{g \in \mathcal{H} \oplus C} \frac{1}{n} \sum_{i=1}^n (y_i - g(\tilde{\mathbf{x}}_i))^2 + \gamma \|g\|_{\mathcal{H}}^2, \quad (6)$$

which we rewrote as:

$$\min_{\beta_0, \boldsymbol{\beta}} \frac{1}{n} \|\mathbf{y} - \beta_0 \mathbf{1} - \mathbf{X}\boldsymbol{\beta}\|^2 + \gamma \|\boldsymbol{\beta}\|^2, \quad (7)$$

In this case,  $\mathcal{H}$  is the RKHS with linear kernel function  $\kappa(\mathbf{x}, \mathbf{x}') = \mathbf{x}^\top \mathbf{x}'$  and  $C = \mathcal{H}_0$  is the linear space of constant functions, which is spanned by the function  $q_1 \equiv 1$ . Moreover,  $\mathbf{K} = \mathbf{X}\mathbf{X}^\top$  and  $\mathbf{Q} = \mathbf{1}$ .

If we appeal to the representer theorem directly, then the problem in (6) becomes, as a result of (3):

$$\min_{\boldsymbol{\alpha}, \eta_0} \frac{1}{n} \|\mathbf{y} - \eta_0 \mathbf{1} - \mathbf{X}\mathbf{X}^\top \boldsymbol{\alpha}\|^2 + \gamma \|\mathbf{X}^\top \boldsymbol{\alpha}\|^2.$$

## Example: Ridge Regression (cont.)

This is a convex optimization problem, and so the solution follows by taking derivatives and setting them to zero. This gives the equations

$$\mathbf{XX}^\top ((\mathbf{XX}^\top + n \gamma \mathbf{I}_n) \boldsymbol{\alpha} + \eta_0 \mathbf{1} - \mathbf{y}) = 0,$$

and

$$n \eta_0 = \mathbf{1}^\top (\mathbf{y} - \mathbf{XX}^\top \boldsymbol{\alpha}).$$

Note that these are equivalent to the equations we found by directly minimizing (7) (assuming that  $n \geq p$  and  $\mathbf{X}$  has full rank  $p$ ).

Equivalently, the solution is found by solving (5):


$$\begin{bmatrix} \mathbf{XX}^\top \mathbf{XX}^\top + n \gamma \mathbf{XX}^\top & \mathbf{XX}^\top \mathbf{1} \\ \mathbf{1}^\top \mathbf{XX}^\top & n \end{bmatrix} \begin{bmatrix} \boldsymbol{\alpha} \\ \eta_0 \end{bmatrix} = \begin{bmatrix} \mathbf{XX}^\top \\ \mathbf{1}^\top \end{bmatrix} \mathbf{y}.$$

This is a system of  $(n + 1)$  linear equations, and is typically of much larger dimension than the  $(p + 1)$  linear equations associated with (7).

## Example: Estimating the Peaks Function

The figure shows the surface plot of the *peaks* function:

$$f(x_1, x_2) = 3(1-x_1)^2 e^{-x_1^2 - (x_2+1)^2} - 10 \left( \frac{x_1}{5} - x_1^3 - x_2^5 \right) e^{-x_1^2 - x_2^2} - \frac{1}{3} e^{-(x_1+1)^2 - x_2^2}.$$



The goal is to learn the function  $y = f(\mathbf{x})$  based on a small set of training data (pairs of  $(\mathbf{x}, y)$  values) indicated by red dots.

## Example: Estimating the Peaks Function

We use the **Gaussian kernel** on  $\mathbb{R}^2$ , and denote by  $\mathcal{H}$  the unique RKHS corresponding to this kernel. We omit the regularization term in (1), and thus our objective is to find the solution to

$$\min_{g \in \mathcal{H}} \frac{1}{n} \sum_{i=1}^n (y_i - g(\mathbf{x}_i))^2.$$


By the representer theorem, the optimal function is of the form

$$g(\mathbf{x}) = \sum_{i=1}^n \alpha_i \exp \left( -\frac{1}{2} \frac{\|\mathbf{x} - \mathbf{x}_i\|^2}{\sigma^2} \right),$$

where  $\boldsymbol{\alpha} := [\alpha_1, \dots, \alpha_n]^\top$  is, by (5), the solution to the set of linear equations  $\mathbf{K}\mathbf{K}^\top \boldsymbol{\alpha} = \mathbf{K}\mathbf{y}$ .

## Example: Estimating the Peaks Function

Note that we are performing regression over the class of functions  $\mathcal{H}$  with an implicit feature space. Due to the representer theorem, the solution to this problem coincides with the solution to the linear regression problem for which the  $i$ -th feature (for  $i = 1, \dots, n$ ) is chosen to be the vector  $[\kappa(\mathbf{x}_1, \mathbf{x}_i), \dots, \kappa(\mathbf{x}_n, \mathbf{x}_i)]^\top$ .



**Figure:** Contour plots for the prediction function  $g$  (left) and the *peaks* function (right).



```
from genham import hammersley
import numpy as np
import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import Axes3D
from matplotlib import cm
from numpy.linalg import norm

import numpy as np
def peaks(x,y):
    z = (3*(1-x)**2 * np.exp(-(x**2) - (y+1)**2)
        - 10*(x/5 - x**3 - y**5) * np.exp(-x**2 - y**2)
        - 1/3 * np.exp(-(x+1)**2 - y**2))
    return(z)

n = 20
x = -3 + 6*hammersley([2,3],n)
z = peaks(x[:,0],x[:,1])
xx, yy = np.mgrid[-3:3:150j,-3:3:150j]
zz = peaks(xx,yy)
plt.contour(xx,yy,zz,levels=50)
fig=plt.figure()
ax = fig.add_subplot(111,projection='3d')
```

```

ax.plot_surface(xx,yy,zz,rstride=1,cstride=1,color='c',alpha=0.3,
               linewidth=0)
ax.scatter(x[:,0],x[:,1],z,color='k',s=20)
plt.show()
sig2 = 0.3 # kernel parameter
def k(x,u):
    return(np.exp(-0.5*norm(x- u)**2/sig2))
K = np.zeros((n,n))
for i in range(n):
    for j in range(n):
        K[i,j] = k(x[i,:],x[j])
alpha = np.linalg.solve(K@K.T, K@z)

N, = xx.flatten().shape
Kx = np.zeros((n,N))
for i in range(n):
    for j in range(N):
        Kx[i,j] = k(x[i,:],np.array([xx.flatten()[j],yy.flatten()[j]]))

g = Kx.T @ alpha
dim = np.sqrt(N).astype(int)
yhat = g.reshape(dim,dim)
plt.contour(xx,yy,yhat,levels=50)

```

# Smoothing Cubic Splines

In the context of data fitting, consider the optimization problem:

$$\min_{g \in \mathcal{G}} \frac{1}{n} \sum_{i=1}^n (y_i - g(x_i))^2 + \gamma \|g''\|^2, \quad (8)$$

where  $\mathcal{G}$  is a suitable function space of twice-differentiable function from  $[0, 1]$  to  $\mathbb{R}$  and  $\|g''\|^2 := \int_0^1 (g''(x))^2 dx$ .

In order to apply the kernel machinery, we want to write this in the form (1), for some RKHS  $\mathcal{H}$  and null space  $\mathcal{H}_0$ .

Clearly, the norm on  $\mathcal{H}$  should be of the form  $\|g\|_{\mathcal{H}} = \|g''\|$  and should be well-defined.

# Smoothing Cubic Splines

This suggests that we take

$$\mathcal{H} = \{g \in L^2[0, 1] : \|g''\| < \infty, g(0) = g'(0) = 0\},$$

with inner product

$$\langle f, g \rangle_{\mathcal{H}} := \int_0^1 f''(x) g''(x) dx.$$

Imposing the condition that  $g(0) = g'(0) = 0$  for functions in  $\mathcal{H}$  will ensure that  $\mathcal{G} = \mathcal{H} \oplus \mathcal{H}_0$  where the null space  $\mathcal{H}_0$  contains only affine functions, as we will see.

# Smoothing Cubic Splines

To see that this  $\mathcal{H}$  is in fact an RKHS, we derive its reproducing kernel. Using integration by parts, write

$$g(x) = \int_0^x g'(s) ds = \int_0^x g''(s) (x-s) ds = \int_0^1 g''(s) (x-s)_+ ds.$$

If  $\kappa$  is a kernel, then by the reproducing property it must hold that

$$g(x) = \langle g, \kappa_x \rangle_{\mathcal{H}} = \int_0^1 g''(s) \kappa_x''(s) ds,$$

so that  $\kappa$  must satisfy  $\frac{\partial^2}{\partial s^2} \kappa(x, s) = (x-s)_+$ , where  $y_+ := \max\{y, 0\}$ .

Therefore, noting that  $\kappa(x, u) = \langle \kappa_x, \kappa_u \rangle_{\mathcal{H}}$ , we have

$$\kappa(x, u) = \int_0^1 \frac{\partial^2 \kappa(x, s)}{\partial s^2} \frac{\partial^2 \kappa(u, s)}{\partial s^2} ds = \frac{\max\{x, u\} \min\{x, u\}^2}{2} - \frac{\min\{x, u\}^3}{6}.$$

# Smoothing Cubic Splines

This is a cubic function with quadratic and cubic terms that misses the constant and linear monomials.

If we now take  $\mathcal{H}_0$  as the space of functions of the following form:

$$h_0 = \eta_1 + \eta_2 x, \quad x \in [0, 1],$$

then (8) is exactly of the form (1).


As a consequence of the representer theorem, the optimal solution to (8) is a **cubic spline** with  $n$  **knots**:

$$g(x) = \eta_1 + \eta_2 x + \sum_{i=1}^n \alpha_i \kappa(x_i, x). \quad (9)$$

The parameters  $\alpha, \eta$  are determined from (3) for instance by solving (5) with matrices  $\mathbf{K} = [\kappa(x_i, x_j)]_{i,j=1}^n$  and  $\mathbf{Q}$  with  $i$ -th row of the form  $[1, x_i]$  for  $i = 1, \dots, n$ .

## Example: Smoothing Spline

Data:  $(0.05, 0.4)$ ,  $(0.2, 0.2)$ ,  $(0.5, 0.6)$ ,  $(0.75, 0.7)$ ,  $(1, 1)$ .



**Figure:** Various cubic smoothing splines for smoothing parameter  $r = 1/(1 + n \gamma) \in \{0.8, 0.99, 0.999, 0.999999\}$ . For  $r = 1$ , the natural spline through the data points is obtained; for  $r = 0$ , the simple linear regression line is found.

```

import matplotlib.pyplot as plt
import numpy as np

x = np.array([[0.05, 0.2, 0.5, 0.75, 1.]]).T
y = np.array([[0.4, 0.2, 0.6, 0.7, 1.]]).T
n = x.shape[0]
r = 0.999
ngamma = (1-r)/r

k = lambda x1, x2 : (1/2)* np.max((x1,x2)) * np.min((x1,x2)) ** 2 \
                    - ((1/6)* np.min((x1,x2))**3)

K = np.zeros((n,n))
for i in range(n):
    for j in range(n):
        K[i,j] = k(x[i], x[j])

Q = np.hstack((np.ones((n,1)), x))
m1 = np.hstack((K @ K.T + (ngamma * K), K @ Q))
m2 = np.hstack((Q.T @ K.T, Q.T @ Q))
M = np.vstack((m1,m2))

c = np.vstack((K, Q.T)) @ y

ad = np.linalg.solve(M,c)

```



```
# plot the curve
```

```
xx = np.arange(0,1+0.01,0.01).reshape(-1,1)
```

```
g = np.zeros_like(xx)
```

```
Qx = np.hstack((np.ones_like(xx), xx))
```

```
g = np.zeros_like(xx)
```

```
N = np.shape(xx)[0]
```

```
Kx = np.zeros((n,N))
```

```
for i in range(n):
```

```
    for j in range(N):
```

```
        Kx[i,j] = k(x[i], xx[j])
```

```
g = g + np.hstack((Kx.T, Qx)) @ ad
```

```
plt.ylim((0,1.15))
```

```
plt.plot(xx, g, label = 'r = {}'.format(r), linewidth = 2)
```

```
plt.plot(x,y, 'b.', markersize=15)
```

```
plt.xlabel('$x$')
```

```
plt.ylabel('$y$')
```

```
plt.legend()
```